

# R - 2ª PARTE

---

## Introducción a la Ciencia de Datos

# Index

- String Manipulation
- Input/Output
- Functions
- R Programming Structures
- Performance Enhancement: Speed and Memory

# String manipulation

- Although R is a statistical language with numeric vectors and matrices playing a central role, character strings are also necessary and R has a number of string-manipulation utilities.
- `nchar()` : finds the length of a string

```
nchar("Ciencia de Datos")
```

```
[1] 16
```

# String manipulation

- `paste()`: concatenates several strings, returning the result in one long string.

```
paste("Ciencia", "de", "Datos")
```

```
[1] "Ciencia de Datos"
```

```
paste("Ciencia", "de", "Datos", sep="_")
```

```
[1] "Ciencia_de_Datos"
```

```
paste("Ciencia", 5+3, "cierta")
```

```
[1] "Ciencia 8 cierta"
```

```
paste(1:3, 1:5, sep="_", collapse="|")
```

```
[1] "1_1|2_2|3_3|1_4|2_5"
```

# String manipulation

- `substr()`: The call `substr(x, start, stop)` returns the substring in the given character position range `start:stop` in the given string `x`.

```
substr("Ciencia de Datos",1,5)
```

```
[1] "Cienc"
```

```
substr("Ciencia de Datos",12)
```

```
[1] "Datos"
```

# String manipulation

- `strsplit()`: The call `strsplit(x, split)` splits a string `x` into an R list of substrings based on another string `split` in `x`.

```
strsplit("2-10-2017", split="-")  
[[1]]  
[1] "2"      "10"     "2017"
```

# Regular expressions

- A regular expression is a kind of wild card.
- It's shorthand to specify broad classes of strings.
- In R, you must pay attention to this point when using the string functions `grep()`, `grepl()`, `regexpr()`, `gregexpr()`, `sub()`, `gsub()`, **and** `strsplit()`.

# Regular expressions

- For example, the expression “[ia]” refers to any string that contains either of the letters *i* or *a*.

```
grep("[ia]",c("Ciencia","de","Datos"))  
[1] 1 3
```

- A period (.) represents any single character.

```
grep(".e",c("Ciencia","de","Datos"))  
[1] 1 2
```



# Regular expressions

- Another example using `strsplit()`:

```
strsplit("a.b.c", ".")
```

```
[[1]]
```

```
[1] "" "" "" "" ""
```

```
strsplit("a.b.c", "[.]")
```

```
[[1]]
```

```
[1] "a" "b" "c"
```

# Input/Output: Keyboard and Monitor

- Suppose we have a file (`file.txt`) with this content:

12

2 5

641

```
scan("file.txt")
```

```
Read 4 items
```

```
[1] 12 2 5 641
```

```
scan("file.txt",what=character())
```

```
Read 4 items
```

```
[1] "12" "2" "5" "641"
```

```
scan("file.txt",sep="\n")
```

```
Read 3 items
```

```
[1] 12 25 641
```

# Input/Output: Keyboard and Monitor

- You can use `scan()` to read from the keyboard by specifying an empty string for the filename:

```
scan("")  
1: 23 4  
3: 2  
4:  
Read 3 items  
[1] 23 4 2
```

- Note that we are prompted with the index of the next item to be input, and we signal the end of input with an empty line.

# Input/Output: Keyboard and Monitor

- If you want to read in a single line from the keyboard use `readline()`:

```
readline("Input data: ")  
Input data: 23 4 2  
[1] "23  4  2"
```

- Note that we are prompted with the index of the next item to be input, and we signal the end of input with an empty line.

# Input/Output: Print to the screen

- `print()` is a *generic* function, so the actual function called will depend on the class of the object that is printed. If, for example, the argument is of class `table`, then the `print.table()` function will be called.

```
x <- 1:3
print(x^2)
[1] 1 4 9
```

# Input/Output: Print to the screen

- It's a little better to use `cat()` instead of `print()`, as the latter can print only one expression and its output is numbered:

```
x <- 1:3
print(x^2)
[1] 1 4 9
cat(x^2)
1 4 9
cat(x^2, x, "hola")
1 4 9 1 2 3 hola
cat(x^2, x, "hola", sep="_")
1_4_9_1_2_3_hola
```

# Input/Output: Reading and Writing files

- We will use of the function `read.table()` to read in a data frame.
- Suppose we have a file `matrix.txt` with the following content:

```
nombre edad
```

```
John 25
```

```
Mary 28
```

```
Jim 19
```

# Input/Output: Reading and Writing files

- The first line contains an optional header, specifying column names. We could read the file this way:

```
read.table("matrix.txt",header=TRUE)
```

	nombre	edad
1	John	25
2	Mary	28
3	Jim	19

- Note that `scan()` would not work here, because our file has a mixture of numeric and character data (and a header).



# Input/Output: Reading and Writing files

- If we want to write a file, we change `read.table()` for `write.table()` function:

```
write.table(matrix(1:6, nrow=2), "output.txt",  
row.names=FALSE, col.names=FALSE)
```

output.txt:

```
1 3 5
```

```
2 4 6
```

# Input/Output: Reading and Writing files

- The function `cat()` can also be used to write to a file, one part at a time:

```
cat("abc\n", file="u.txt")
```

```
cat("de\n", file="u.txt", append=TRUE)
```

```
u.txt:
```

```
abc
```

```
de
```

# Functions

- A *function* is a group of instructions that takes inputs, uses them to compute other values, and returns a result.

```
suma <- function(x,y) {  
  x + y  
}  
  
suma(2,3)  
[1] 5
```

# Functions

- Default values:

```
suma <- function(x, y=5) {  
  x + y  
}  
  
suma(2)  
[1] 7
```

# Functions

- Default values:

```
suma <- function(x=5,y) {  
  x + y  
}
```

```
suma(y=2)
```

```
[1] 7
```

# Programming Structures

- R is a block-structured language (like C, C++, Python, Perl, and so on).
- Blocks are delineated by braces, while statements are separated by newline characters or, optionally, by semicolons.
- As with many scripting languages, we do not “declare” variables in R, therefore we have to take care of possible variable scoping issues.

# Basic R operators

Operation	Description
<code>x + y</code>	Addition
<code>x - y</code>	Subtraction
<code>x * y</code>	Multiplication
<code>x / y</code>	Division
<code>x ^ y</code>	Exponentiation
<code>x %% y</code>	Modular arithmetic
<code>x %/% y</code>	Integer division
<code>x == y</code>	Test for equality
<code>x &lt;= y</code>	Test for less than or equal to
<code>x &gt;= y</code>	Test for greater than or equal to
<code>x &amp;&amp; y</code>	Boolean AND for scalars
<code>x    y</code>	Boolean OR for scalars
<code>x &amp; y</code>	Boolean AND for vectors (vector x,y,result)
<code>x   y</code>	Boolean OR for vectors (vector x,y,result)
<code>!x</code>	Boolean negation

# Control Statements: if-else

- The syntax for if-else looks like:

```
if (r == 4) {  
    x <- 1  
}  
else {  
    x <- 3  
    y <- 4  
}
```



# Control Statements: if-else

- An if-else statement works as a function call, and as such, it returns the last value assigned.

```
if (x == 2) y <- x else y <- x+1
```

```
y <- if (x == 2) x else x+1
```

# Control Statements: if-else

- When working with vectors, use the `ifelse()` function.
- The form is: `ifelse(b, u, v)` where **b** is a Boolean vector, and **u** and **v** are vectors.
- The return value is itself a vector: element *i* is `u[i]` if `b[i]` is true, or `v[i]` if `b[i]` is false.

# Control Statements: if-else

```
x <- 1:10
ifelse(x %% 2 == 0, "par", "impar")

[1] "impar" "par" "impar" "par" "impar"
"par" "impar" "par" "impar" "par"
```

```
x <- c(5, 2, 9, 12)
ifelse(x > 6, 2*x, 3*x)

[1] 15 6 18 24
```

# Control Statements: loops

- One of the major themes of R programming is to avoid loops if possible; if not, keep loops simple.
- We have:
  - For loops
  - While loops
  - Repeat loops

# Control Statements: loops

- For loops:

```
k <- 0  
for(n in x) {  
  if (n %% 2 == 1) k <- k+1  
}  
k  
[1] 5
```

# Control Statements: loops

- For loops:

```
k <- 0
for (i in 1:length(x)) {
  if (x[i] %% 2 == 1) k <- k+1
}
k
[1] 5
```

# Control Statements: loops

- For loops:

```
for (fn in c("file1", "file2")) print(scan(fn))
```

```
Read 6 items
```

```
[1] 1 2 3 4 5 6
```

```
Read 3 items
```

```
[1] 5 12 13
```

# Control Statements: While loops

- While loops:

```
i <- 1  
while (i <= 10) i <- i+4  
i  
[1] 13
```



# Control Statements: Repeat loops

- Repeat loops:

```
i <- 1
repeat {
  i <- i+4
  if (i > 10) break
}
i
```

[1] 13

# Looping, the R way...

- In R you have more options when doing repeating calculations:

function	input	output	comment
apply	matrix or array	vector or array or list	
lapply	list or vector	list	
sapply	list or vector	vector or matrix or list	simplify
vapply	list or vector	vector or matrix or list	safer simplify
tapply	data, categories	array or list	ragged
mapply	lists and/or vectors	vector or matrix or list	multiple
rapply	list	vector or list	recursive
eapply	environment	list	
dendrapply	dendogram	dendogram	
rollapply	data	similar to input	package zoo

# The apply function

- This is the general form of apply for matrices:

```
apply(m, dimcode, f, fargs)
```

- where:
  - *m* is the matrix.
  - *dimcode* is the dimension, equal to 1 if the function applies to rows or 2 for columns.
  - *f* is the function to be applied.
  - *fargs* is an optional set of arguments to be supplied to *f*.

# The apply function

- Apply over the margins of an array (e.g., the rows or columns of a matrix).

```
z <- matrix(c(1,1,2,2,3,3), nrow=3, byrow=TRUE)
```

```
apply(z,1,mean)
```

```
[1] 1 2 3
```

```
apply(z,2,mean)
```

```
[1] 2 2
```

```
apply(z,1:2,mean)
```

```
      [,1] [,2]
```

```
[1,]     1     1
```

```
[2,]     2     2
```

```
[3,]     3     3
```

# The lapply function

- Returns a list of the same length as the input data `x`, each element of which is the result of applying a function to the corresponding element of `x`.

```
z <- list(vector=c(1,2,3), matriz=matrix(1,nrow=2,ncol=2))
```

```
lapply(z,mean)
```

```
$vector
```

```
[1] 2
```

```
$matriz
```

```
[1] 1
```

# The supply function

- Apply over an object and return a simplified object (an array) if possible.

```
supply(z, mean)
```

```
vector matriz
```

```
      2      1
```

## Performance Enhancement: Speed and Memory

- In order to have a fast-running program, you may need to use more memory space.
- On the other hand, in order to conserve memory space, you might need to settle for slower code.
- R is an interpreted language.
  - Many of the commands are written in C and thus do run in fast machine code. But other commands, and your own R code, are pure R and thus interpreted.
- All objects in an R session are stored in memory.
  - More precisely, all objects are stored in R's memory address space.

# Performance Enhancement: Tips

- Optimize your R code through **vectorization**, use of byte-code compilation, and other approaches.
- Write the key, CPU-intensive parts of your code in a compiled language such as C/C++.
- Write your code in some form of parallel R.



# Vectorization

- Loops:
  - It's important to understand that simply rewriting code to avoid loops will not necessarily make the code faster.
  - However, in some cases, dramatic speedup may be attained, usually through vectorization.

# Vectorization

```
for (i in 1:length(x)) z[i] <- x[i] + y[i]
```

**VS.**

```
z <- x + y
```

```
x <- runif(1000000)
y <- runif(1000000)
z <- vector(length=1000000)
system.time(for (i in 1:length(x)) z[i] <- x[i] + y[i])
  user  system elapsed
1.504    0.102    1.724
system.time(z <- x + y)
  user  system elapsed
0.002    0.004    0.007
```

# Vectorization

- Examples of other vectorized functions that may speed up your code are `ifelse()`, `which()`, `where()`, `any()`, `all()`, `cumsum()`, and `cumprod()`.
- In the matrix case, you can use `rowSums()`, `colSums()`, and so on.
- In “all possible combinations” types of settings, `combin()`, `outer()`, `lower.tri()`, `upper.tri()`, or `expand.grid()` may be just what you need.
- Though `apply()` eliminates an explicit loop, it is actually implemented in R rather than C and thus will usually not speed up your code. However, the other apply functions, such as `lapply()`, can be very helpful in speeding up your code.

# Performance Enhancement

- Slow algorithm in R

```
xs <- runif(1000)
res <- c()
for (x in xs) {
  # This is slow!
  res <- c(res, sqrt(x))
}
```

# Performance Enhancement

- Faster algorithm in R

```
xs <- runif(1000)
res <- numeric(length(xs))
for (i in seq_along(xs)) {
  res[i] <- sqrt(xs[i])
}
```

# Performance Enhancement

- Slower algorithm in R

```
amat <- matrix(1:20, nrow=4)
bmat <- matrix(NA, nrow(amat)/2, ncol(amat))
for(i in 1:nrow(bmat))
  bmat[i,] <- amat[2*i-1,] * amat[2*i,]
```

# Performance Enhancement

- Faster algorithm in R

```
amat <- matrix(1:20, nrow=4)
```

```
bmat2 <- amat[seq(1, nrow(amat), by=2),] *  
amat[seq(2, nrow(amat), by=2),]
```

# Over-Vectorizing

- It is a good thing to want to vectorize when there is no effective way to do so. It is a bad thing to attempt it anyway.
- A common reflex is to use a function in the apply family. This is not vectorization, it is loop-hiding.
  - The apply function has a for loop in its definition.
  - Use an explicit for loop when each iteration is a non-trivial task. But a simple loop can be more clearly and compactly expressed using an apply function.



# Over-Vectorizing

- Suppose that we want to find all of the sets of three positive integers that sum to 6, where the order matters:

```
the.seq <- 1:4  
which(outer(outer(the.seq, the.seq, '+'),  
the.seq, '+') == 6, arr.ind=TRUE)
```

# Performance Enhancement

- Some things are not possible to vectorize.
- If you need to use a loop, then:
  - Put as much outside of loops as possible. Make the number of iterations as small as possible.

# References

- Norman Matloff. 2011. The Art of R Programming: A Tour of Statistical Software Design (1st ed.). No Starch Press, San Francisco, CA, USA.
- Patrick Burns. 2011. The R Inferno.